

# Classic Classifier

```
In [ ]: %env CUDA_VISIBLE_DEVICES=1
```

env: CUDA\_VISIBLE\_DEVICES=1

```
In [ ]: import numpy as np
import cudf
import cuml
from cupy import asnumpy
from joblib import dump, load
```

```
In [ ]: def calc_f1(cm):
    # Extracting True Positives, False Positives, False Negatives
    TP = cm[0][0]
    FP = cm[0][1]
    FN = cm[1][0]
    # TN = confusion_matrix[1][1] # True Negatives are not needed for F1

    # Calculating Precision and Recall
    precision = TP / (TP + FP) if (TP + FP) > 0 else 0
    recall = TP / (TP + FN) if (TP + FN) > 0 else 0

    # Calculating F1 Score
    if precision + recall == 0: # Avoiding division by zero
        f1_score = 0
    else:
        f1_score = 2 * (precision * recall) / (precision + recall)

    return f1_score
```

## Load Data

```
In [ ]: train_data_path = "./data/train.csv"
```

```
In [ ]: df = cudf.read_csv(train_data_path)
df.describe()
```

```
/home/dx/miniconda3/envs/rapids-24.02/lib/python3.10/site-packages/cudf/core/dataframe.py:5106: FutureWarning: `datetime_is_numeric` is deprecated.
Specify `datetime_is_numeric=True` to silence this warning and adopt the future behavior now.
  warnings.warn(
/home/dx/miniconda3/envs/rapids-24.02/lib/python3.10/site-packages/cudf/core/series.py:3319: FutureWarning: `datetime_is_numeric` is deprecated and
will be removed in a future release. Specify `datetime_is_numeric=True` to
silence this warning and adopt the future behavior now.
  warnings.warn(
```

Out [ ]:

|              | ind_recommended | activation   | customer_digital_activity_04 | customer_sj |
|--------------|-----------------|--------------|------------------------------|-------------|
| <b>count</b> | 1.222998e+07    | 1.222998e+07 | 1.472619e+06                 | 1.018       |
| <b>mean</b>  | 1.264980e-01    | 5.725000e-03 | 9.745803e+00                 | 1.342       |
| <b>std</b>   | 3.324100e-01    | 7.544700e-02 | 4.156839e+01                 | 6.454       |
| <b>min</b>   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 1.0000      |
| <b>25%</b>   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 3.844       |
| <b>50%</b>   | 0.000000e+00    | 0.000000e+00 | 2.000000e+00                 | 6.437       |
| <b>75%</b>   | 0.000000e+00    | 0.000000e+00 | 6.000000e+00                 | 1.138       |
| <b>max</b>   | 1.000000e+00    | 1.000000e+00 | 8.560000e+02                 | 2.489       |

8 rows x 71 columns

## Data Cleanup

In [ ]: `df.drop(["customer", "merchant"], axis=1, inplace=True)`  
`df.describe()`

```
/home/dx/miniconda3/envs/rapids-24.02/lib/python3.10/site-packages/cudf/core/dataframe.py:5106: FutureWarning: `datetime_is_numeric` is deprecated.
Specify `datetime_is_numeric=True` to silence this warning and adopt the future
behavior now.
  warnings.warn(
/home/dx/miniconda3/envs/rapids-24.02/lib/python3.10/site-packages/cudf/core/series.py:3319: FutureWarning: `datetime_is_numeric` is deprecated and
will be removed in a future release. Specify `datetime_is_numeric=True` to
silence this warning and adopt the future behavior now.
  warnings.warn(
```

Out [ ]:

|              | ind_recommended | activation   | customer_digital_activity_04 | customer_sj |
|--------------|-----------------|--------------|------------------------------|-------------|
| <b>count</b> | 1.222998e+07    | 1.222998e+07 | 1.472619e+06                 | 1.018       |
| <b>mean</b>  | 1.264980e-01    | 5.725000e-03 | 9.745803e+00                 | 1.342       |
| <b>std</b>   | 3.324100e-01    | 7.544700e-02 | 4.156839e+01                 | 6.454       |
| <b>min</b>   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 1.0000      |
| <b>25%</b>   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 3.844       |
| <b>50%</b>   | 0.000000e+00    | 0.000000e+00 | 2.000000e+00                 | 6.437       |
| <b>75%</b>   | 0.000000e+00    | 0.000000e+00 | 6.000000e+00                 | 1.138       |
| <b>max</b>   | 1.000000e+00    | 1.000000e+00 | 8.560000e+02                 | 2.489       |

8 rows x 69 columns

In [ ]: `print(df.shape)`

(12229978, 69)

```
In [ ]: threshold = len(df.columns) * 0.8
df_cleaned = df.dropna(thresh=threshold)
df_cleaned.shape
```

```
Out[ ]: (1373560, 69)
```

```
In [ ]: df_cleaned.describe()
```

```
/home/dx/miniconda3/envs/rapids-24.02/lib/python3.10/site-packages/cudf/core/dataframe.py:5106: FutureWarning: `datetime_is_numeric` is deprecated. Specify `datetime_is_numeric=True` to silence this warning and adopt the future behavior now.
  warnings.warn(
/home/dx/miniconda3/envs/rapids-24.02/lib/python3.10/site-packages/cudf/core/series.py:3319: FutureWarning: `datetime_is_numeric` is deprecated and will be removed in a future release. Specify `datetime_is_numeric=True` to silence this warning and adopt the future behavior now.
  warnings.warn(
```

```
Out[ ]:
```

|       | ind_recommended | activation   | customer_digital_activity_04 | customer_s   |
|-------|-----------------|--------------|------------------------------|--------------|
| count | 1.373560e+06    | 1.373560e+06 | 1.373560e+06                 | 1.373560e+06 |
| mean  | 2.369470e-01    | 1.963200e-02 | 3.283215e+00                 | 1.251000e+00 |
| std   | 4.252100e-01    | 1.387330e-01 | 2.212771e+01                 | 4.676000e+00 |
| min   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 1.000000e+00 |
| 25%   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 4.373000e+00 |
| 50%   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 6.745000e+00 |
| 75%   | 0.000000e+00    | 0.000000e+00 | 1.000000e+00                 | 1.129000e+01 |
| max   | 1.000000e+00    | 1.000000e+00 | 8.470000e+02                 | 1.500000e+01 |

8 rows x 69 columns

```
In [ ]: are_any_nulls = df_cleaned.isnull().any().any()
are_any_nulls
```

```
Out[ ]: True
```

```
In [ ]: for column in df.columns:
    most_frequent = df_cleaned[column].value_counts().index[0] # Get the most frequent value
    df_cleaned[column] = df_cleaned[column].fillna(most_frequent) # Fill with the most frequent value
```

```
In [ ]: are_any_nulls = df_cleaned.isnull().any().any()
are_any_nulls
```

```
Out[ ]: False
```

```
In [ ]: from cuml.model_selection import train_test_split

X = df_cleaned.drop(["activation", "ind_recommended"], axis=1)
y = df_cleaned["activation"]
print(X.shape, y.shape)
```

```
(1373560, 67) (1373560,)
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

## Random Forest

```
In [ ]: from cuml.ensemble import RandomForestClassifier
```

```
rfc = RandomForestClassifier(n_estimators=100,  
                             max_depth=16)
```

```
In [ ]: rfc.fit(X_train, y_train)
```

```
Out[ ]: ▼ RandomForestClassifier ⓘ  
RandomForestClassifier()
```

```
In [ ]: from cuml.metrics import accuracy_score
```

```
y_train_pred = rfc.predict(X_train)  
rf_train_accuracy = accuracy_score(y_train, y_train_pred)  
print("Train accuracy: ", rf_train_accuracy)
```

```
y_test_pred = rfc.predict(X_test)  
rf_test_accuracy = accuracy_score(y_test, y_test_pred)  
print("Test accuracy: ", rf_test_accuracy)
```

Train accuracy: 0.9867188334465027

Test accuracy: 0.9827929139137268

```
In [ ]: from cuml.metrics import confusion_matrix
```

```
print("Train: ", confusion_matrix(y_train, y_train_pred, convert_dtype=True))  
print("Test: ", confusion_matrix(y_test, y_test_pred, convert_dtype=True))
```

Train: [[1077149 0]

[ 14594 7105]]

Test: [[269254 191]

[ 4536 731]]

```
In [ ]: print("F1: ", calc_f1(confusion_matrix(y_test, y_test_pred, convert_dtype
```

F1: 0.991289301229659

```
In [ ]: dump(rfc, 'rfc.joblib')
```

```
Out[ ]: ['rfc.joblib']
```

```
In [ ]:
```

## SVM

```
In [ ]: from cuml.svm import SVC
```

```
svm = SVC(kernel='rbf', class_weight='balanced')
```

```
In [ ]: svm.fit(X_train, y_train)
```

```
Out[ ]: SVC
SVC()
```

```
In [ ]: from cuml.metrics import accuracy_score

y_train_pred = svm.predict(X_train)
svm_train_accuracy = accuracy_score(y_train, y_train_pred)
print("Train accuracy: ", svm_train_accuracy)

y_test_pred = svm.predict(X_test)
svm_test_accuracy = accuracy_score(y_test, y_test_pred)
print("Test accuracy: ", svm_test_accuracy)
```

```
Train accuracy: 0.7885931730270386
Test accuracy: 0.7888770699501038
```

```
In [ ]: from cuml.metrics import confusion_matrix

print("Train: ", confusion_matrix(y_train, y_train_pred, convert_dtype=True))
print("Test: ", confusion_matrix(y_test, y_test_pred, convert_dtype=True))
```

```
Train: [[853296 223853]
 [ 8451 13248]]
Test: [[213519 55926]
 [ 2072 3195]]
```

```
In [ ]: print("F1: ", calc_f1(confusion_matrix(y_test, y_test_pred, convert_dtype=True)))

F1: 0.8804253704879638
```

```
In [ ]: dump(svm, 'svm.joblib')
```

```
Out[ ]: ['svm.joblib']
```